

```
In [47]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
In [35]: # abrimos nuestro archivo con el que trabajaremos

emisiones = pd.read_excel('C:/Users/Isaac/Desktop/IHD/EBAC DT/M20 DS/FuelConsumptionCo2.xlsx')
emisiones.head()
```

```
Out[35]:
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION	FUELTYPE	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY
0	2022	Acura	ILX	Compact	2.4	4	AM8	Z	9.9	7.0
1	2022	Acura	MDX SH-AWD	SUV: Small	3.5	6	AS10	Z	12.6	9.4
2	2022	Acura	RDX SH-AWD	SUV: Small	2.0	4	AS10	Z	11.0	8.6
3	2022	Acura	RDX SH-AWD A-SPEC	SUV: Small	2.0	4	AS10	Z	11.3	9.1

```
In [36]: emisiones.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 945 entries, 0 to 944
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   MODELYEAR                            945 non-null    int64
1   MAKE                                945 non-null    object
2   MODEL                                945 non-null    object
3   VEHICLECLASS                         945 non-null    object
4   ENGINE SIZE                          945 non-null    float64
5   CYLINDERS                           945 non-null    int64
6   TRANSMISSION                        945 non-null    object
7   FUELTYPE                             945 non-null    object
8   FUELCONSUMPTION_CITY                 945 non-null    float64
9   FUELCONSUMPTION_HWY                 945 non-null    float64
10  FUELCONSUMPTION_COMB                 945 non-null    float64
11  FUELCONSUMPTION_COMB_MPG             945 non-null    int64
12  CO2EMISSIONS                        945 non-null    int64
dtypes: float64(4), int64(4), object(5)
memory usage: 96.1+ KB
```

```
In [37]: # creamos un grafico de calor para verificar si hay datos faltantes

plt.figure(figsize = (10, 5))
sns.heatmap(emisiones.isnull(), yticklabels = False, cmap = 'crest')

# determinamos que no hay valores faltantes
```

```
Out[37]: <Axes: >
```



In [38]: `# eliminamos variables categoricas`

```
emisiones.drop(['MAKE', 'MODEL', 'VEHICLECLASS', 'TRANSMISSION', 'FUELTYPE'], axis = 1, inplace = True)
emisiones
```

Out[38]:

ENGINE_SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY	FUELCONSUMPTION_COMB	FUELCONSUMPTION_COMB_MPG	CO2EMISSIONS
2.4	4	9.9	7.0	8.6	33	200
3.5	6	12.8	9.4	11.2	25	263
2.0	4	11.0	8.6	9.9	29	232
2.0	4	11.3	9.1	10.3	27	242
2.0	4	11.2	8.0	9.8	29	230
...
2.0	4	10.7	7.7	9.4	30	219
2.0	4	10.5	8.1	9.4	30	219
2.0	4	11.0	8.7	9.9	29	232
2.0	4	11.5	8.4	10.1	28	236
2.0	4	12.4	8.9	10.8	28	252

nns

In [11]: `emisiones.shape`

Out[11]: (946, 8)

In [39]: `# definimos los valores para 'X', 'y'`

```
X = emisiones.drop('CO2EMISSIONS', axis = 1)
y = emisiones['CO2EMISSIONS']
```

In [40]: `X`

Out[40]:

	MODELYEAR	ENGINE_SIZE	CYLINDERS	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY	FUELCONSUMPTION_COMB	FUELCONSUMPTION_COMB_MPG
0	2022	2.4	4	9.9	7.0	8.6	
1	2022	3.5	6	12.8	9.4	11.2	
2	2022	2.0	4	11.0	8.6	9.9	
3	2022	2.0	4	11.3	9.1	10.3	
4	2022	2.0	4	11.2	8.0	9.8	
...
940	2022	2.0	4	10.7	7.7	9.4	
941	2022	2.0	4	10.5	8.1	9.4	
942	2022	2.0	4	11.0	8.7	9.9	
943	2022	2.0	4	11.5	8.4	10.1	
944	2022	2.0	4	12.4	8.9	10.8	

945 rows x 7 columns

In [41]: `y`

Out[41]:

0	200
1	263
2	232
3	242
4	230
...	...
940	219
941	219
942	232
943	236
944	252

Name: CO2EMISSIONS, Length: 945, dtype: int64

```
In [50]: # Dividimos en grupos de entrenamiento(training) y prueba(test)

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state = 1)
```

CONSTRUCCION DE MODELO DE REGRESION LINEAL MULTIPLE

```
In [51]: from sklearn.linear_model import LinearRegression

linreg = LinearRegression()
```

```
In [52]: # entrenamos el modelo

linreg.fit(X_train, y_train)
```

```
Out[52]: LinearRegression
LinearRegression()
```

COEFICIENTES DE REGRESION

```
In [54]: print('Intercepto: ', linreg.intercept_)
print('Coeficientes: ', linreg.coef_)

Intercepto: 124.05613100622463
Coeficientes: [ 0.          -0.01515148  6.43036425 -1.32628678  0.52702525 14.53834062
 -1.88292508]
```

```
In [55]: # realizamos predicciones (probamos modelo)

linregpred = linreg.predict(X_test)
linregpred
```

```
Out[55]: array([318.82062694, 322.84703822, 345.23551728, 174.78239603,
196.73994807, 250.37345085, 240.2772048 , 332.29424162,
327.80185548, 173.64652185, 275.75871276, 227.67606101,
323.15858211, 264.7147751 , 175.92922152, 267.50620255,
162.5548175 , 220.68619289, 337.80918664, 326.56057625,
322.84703822, 339.48137878, 259.20584163, 220.68619289,
339.76297512, 201.8255166 , 364.62108712, 334.53095579,
286.6535063 , 137.95706661, 179.28808399, 175.92922152,
238.60801519, 291.97541448, 312.36078823, 291.32632034,
330.50151422, 317.77510602, 327.80185548, 218.69601311,
302.19635808, 294.44143376, 228.58586706, 182.33789374,
170.20435766, 247.91950372, 193.35806215, 343.30481828,
152.91430008, 184.32703834, 291.68467823, 239.45777535,
170.44542168, 262.19926692, 112.39018005, 312.507081 ,
343.13766885, 335.91261837, 307.08026329, 385.77737538,
198.41216797, 150.50258803, 186.61208853, 304.78932457,
```

AJUSTAMOS EL MODELO CON METRICAS

```
In [57]: from sklearn.metrics import r2_score
from sklearn import metrics
```

```
In [60]: # probamos las metricas

print('Valor de R2: ', r2_score(y_test, linregpred))
print('Error Absoluto Medio: ', metrics.mean_absolute_error(y_test, linregpred))
print('Error Cuadrático Medio: ', metrics.mean_squared_error(y_test, linregpred))
print('Raíz del Error Cuadrático Medio: ', np.sqrt(metrics.mean_squared_error(y_test, linregpred)))

Valor de R2: 0.9775327964689398
Error Absoluto Medio: 6.4183220979061755
Error Cuadrático Medio: 85.97881014782169
Raíz del Error Cuadrático Medio: 9.272475944849988
```

REGRESION RIDGE

```
In [61]: from sklearn.linear_model import Ridge
```

```
In [68]: # probamos con un valor de alpha = 1

ridgereg = Ridge(alpha = 0.1)

# entrenamos el modelo

ridgereg.fit(X_train, y_train)
```

```
Out[68]:
+   Ridge
+   Ridge(alpha=0.1)
```

```
In [69]: # realizamos predicciones

ridgeregpred = ridgereg.predict(X_test)
```

```
In [70]: # probamos con las metricas de ajuste

print('Valor de R2: ', r2_score(y_test, ridgeregpred))
print('Error Absoluto Medio: ', metrics.mean_absolute_error(y_test, ridgeregpred))
print('Error Cuadrático Medio: ', metrics.mean_squared_error(y_test, ridgeregpred))
print('Raíz del Error Cuadrático Medio: ', np.sqrt(metrics.mean_squared_error(y_test, ridgeregpred)))

Valor de R2: 0.9775405388157286
Error Absoluto Medio: 6.415135896250665
Error Cuadrático Medio: 85.94918128173981
Raíz del Error Cuadrático Medio: 9.27087812894441
```

COEFICIENTES DEL MODELO RIDGE

```
In [64]: print('Intercepto: ', ridgereg.intercept_)
print('Coeficientes: ', ridgereg.coef_)

Intercepto: 123.99617710866374
Coeficientes: [ 0.          -0.017015    6.42593926 -0.50568931  1.19661489 13.05363848
-1.88187263]
```

```
In [71]: # buscamos valor optimo de alpha

alpha_range = 10. ** np.arange(-2, 3)
alpha_range
```

```
Out[71]: array([1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+02])
```

```
In [72]: # utilizamos CV para encontrar la combinacion del valor optimo

from sklearn.linear_model import RidgeCV
```

```
In [73]: ridgeregcv = RidgeCV(alphas = alpha_range)
```

```
In [74]: # entrenamos el modelo

ridgeregcv.fit(X_train, y_train)
ridgeregcv.alpha_
```

```
Out[74]: 100.0
```

```
In [75]: # realizamos la predicción utilizando el mejor valor de alpha

ridgeregcvpred = ridgeregcv.predict(X_test)
```

```
In [76]: # probamos con las metricas de ajuste

print('Valor de R2: ', r2_score(y_test, ridgeregcvpred))
print('Error Absoluto Medio: ', metrics.mean_absolute_error(y_test, ridgeregcvpred))
print('Error Cuadrático Medio: ', metrics.mean_squared_error(y_test, ridgeregcvpred))
print('Raíz del Error Cuadrático Medio: ', np.sqrt(metrics.mean_squared_error(y_test, ridgeregcvpred)))

Valor de R2: 0.9779646398338884
Error Absoluto Medio: 6.348640227898685
Error Cuadrático Medio: 84.32620667017098
Raíz del Error Cuadrático Medio: 9.182930178879232
```

COEFICIENTES DE REGRESION DEL MODELO RIDGE USANDO CV

```
In [77]: print('Intercepto: ', ridgeregcv.intercept_)
print('Coeficientes: ', ridgeregcv.coef_)

Intercepto: 129.98702013615804
Coeficientes: [ 0.          1.15258677  5.30842335  4.44225152  4.49815493  4.55653156
-1.95117901]
```

REGRESIÓN LASSO

```
In [78]: # prueba con valor de alpha = 0.001
from sklearn.linear_model import Lasso
```

```
In [79]: lassoreg = Lasso(alpha = 0.001)
```

```
In [80]: # entrenamos el modelo
lassoreg.fit(X_train, y_train)
```

```
Out[80]: Lasso
Lasso(alpha=0.001)
```

COEFICIENTES DE REGRESION DEL MODELO LASSO

```
In [81]: print('Intercepto: ', lassoreg.intercept_)
print('Coeficientes: ', lassoreg.coef_)

Intercepto: 123.13486003787818
Coeficientes: [ 0.          -0.05568687  6.39362175  7.1821322   7.48383257 -0.85056791
 -1.86749083]
```

```
In [82]: # prueba con alpha = 0.01
lassoreg = Lasso(alpha = 0.01)
```

```
In [83]: # entrenamos el modelo
lassoreg.fit(X_train, y_train)
```

```
Out[83]: Lasso
Lasso(alpha=0.01)
```

```
In [84]: # imprimimos coeficientes

print('Intercepto: ', lassoreg.intercept_)
print('Coeficientes: ', lassoreg.coef_)

Intercepto: 123.34866942884304
Coeficientes: [ 0.          -0.          6.36561443  6.29526617  6.74930805  0.75481726
 -1.87010303]
```

```
In [85]: # realizamos las predicciones
lassoregpred = lassoreg.predict(X_test)
```

```
In [86]: print('Valor de R2: ', r2_score(y_test, lassoregpred))
print('Error Absoluto Medio: ', metrics.mean_absolute_error(y_test, lassoregpred))
print('Error Cuadrático Medio: ', metrics.mean_squared_error(y_test, lassoregpred))
print('Raíz del Error Cuadrático Medio: ', np.sqrt(metrics.mean_squared_error(y_test, lassoregpred)))

Valor de R2: 0.977565671226078
Error Absoluto Medio: 6.4022959826076935
Error Cuadrático Medio: 85.85300310206533
Raíz del Error Cuadrático Medio: 9.265689564304717
```

SELECCION DEL VALOR OPTIMO PARA EL MODELO LASSO

```
In [87]: from sklearn.linear_model import LassoCV
```

```
In [88]: lassoregcv = LassoCV(n_alphas = 100, random_state = 1)
```

```
In [89]: # entrenamos el modelo
lassoregcv.fit(X_train, y_train)
```

```
Out[89]: LassoCV
LassoCV(random_state=1)
```

```
In [90]: print('Alpha Optimo: ', lassoregcv.alpha_)

Alpha Optimo: 0.45442206714715
```

COEFICIENTES DE REGRESION DE LASSOCV

```
In [91]: print('Intercepto: ', lassoregcv.intercept_)
print('Coeficientes: ', lassoregcv.coef_)

Intercepto: 126.10535289477207
Coeficientes: [ 0.          0.          6.14099602  5.4462973   5.62023828  2.61988286
 -1.90585504]
```

```
In [92]: # prediccion de lasso con el valor optimo
lassoregcvpred = lassoregcv.predict(X_test)
```

```
In [93]: print('Valor de R2', r2_score(y_test, lassoregcvpred))
print('Error Absoluto Medio', metrics.mean_absolute_error(y_test, lassoregcvpred))
print('Error Cuadratico Medio', metrics.mean_squared_error(y_test, lassoregcvpred))
print('Raiz del Error Caudratico Medio', np.sqrt(metrics.mean_squared_error(y_test, lassoregcvpred)))

Valor de R2 0.9775539332111336
Error Absoluto Medio 6.405732289229923
Error Cuadratico Medio 85.89792282503036
Raiz del Error Caudratico Medio 9.268113228971167
```

RESULTADOS

```
In [95]: print('Regresion Lienal Multiple: %f' % r2_score(y_test, linregpred))
print('Regresion Ridge: %f' % r2_score(y_test, ridgeregcvpred))
print('Regresion Lasso: %f' % r2_score(y_test, lassoregcvpred))

Regresion Lienal Multiple: 0.977533
Regresion Ridge: 0.977965
Regresion Lasso: 0.977554
```

CONCLUSIÓN.

En este caso los 3 modelos dieron excelentes resultados.

El modelo que mas destaco fue la Regresion Lasso obteniendo un valor de 0.9779 de los 3 modelos es el mas optimo ya que se acerca al valor de 1.

In []:

